

Tobias Treppmann

Agentic AI Specialist — a production AI agent for a customer-service team, and the knowledge source behind it

LOCATION Berlin, Germany EMAIL hello@treppmann.design PORTFOLIO treppmann.design/for/docmorris LINKEDIN linkedin.com/in/tobiastreppmann

PROFILE

I built the evaluation panel that judges an AI pipeline's output, and shipped an AI agent for a customer-service team. Nine criteria, each judged separately, any failure rejecting the draft — no golden set. I wrote the agent's system prompt and author the knowledge source that decides what it can say. Designed and directed; implemented with AI coding agents.

EXPERIENCE

Head of Product Design

2025 – Present

Bibel TV (broadcaster in linear-to-streaming transformation) Hamburg + Berlin, Germany

- Shipped an AI support agent into production for the viewer-services team — wrote the system prompt, an explicit restricted-topic list, and the scope rule that refuses to draft outside a viewer enquiry.
- Author and structure the knowledge source it is grounded on — a web-hosted handbook summary. That page decides the agent's answerable scope, and is where most of the work actually goes.
- Defined what a good draft looks like per message type with the support team before tuning anything, then tested and optimized against that standard.
- Rollout was voluntary — a self-test week, then one-to-ones. The skeptics came round after their hardest case, not the demo; every send stays a person's decision.
- Built an AI fundraising system end to end — analysis, prompt architecture, prototype, and a production Claude skill across in-app, web and video, with follow-up-question loops.
- Built a token suggestion engine, then constrained it after it offered a spacing token as a font-size fix on a numeric coincidence — now bounded to the correct property family, returning nothing rather than a plausible wrong answer.
- Enforcement runs as six automated checks, identical locally and in CI, refusing an unresolvable token reference or a hardcoded value; prototypes are built in production code against real APIs, with an external agency shipping from it.
- No in-house design team, so I own UI quality and consistency across iOS, Android, web and connected TV, with nobody to delegate it to. Product managers and product owners build day to day, and an external engineering agency ships from my prototype code across a company boundary.

Founder · Aperto

2025 – Present

Aperto (aperto.bible) — open Bible translation tooling and a multilingual content pipeline Berlin, Germany · independent

- Built a multi-layer evaluation panel for LLM output. Nine criteria, each judged in its own model call with no shared context, binary YES or NO, any single failure rejecting the candidate.
- No golden set, and that was a decision, not an omission. A reference key existed; removed in July 2026 after it hid eight of eleven failures on a chapter it had passed.
- Criteria are tiered and scoped — a floor tier for correctness and mechanics, a primary tier for register and reader impact — and judge reasons go back to the generator verbatim inside a closed revision loop.
- Deterministic and statistical judges sit beside the model ones — style-distance against a reference corpus, documented as a floor not a verdict, plus drift detection against published translations and a judge-drift calibration control.
- Prompt changes validated by pre-registered A/B probes with byte-identical control arms, plus blind preference judging against the current baseline.
- Built the retrieval layer deliberately without embeddings — commentaries chunked by pericope into chapter-keyed JSON, so an exegesis step pulls exactly the scholarship for its passage.
- 139 loop runs and 3,138 individual criterion verdicts since July 2026, across 64 German, 55 English and 42 Polish chapters.
- Designed and directed across 1,384 commits; implemented with AI coding agents.

Head of Platform and Design

2024 – 2025

ChurchDesk (vertical ERP for churches and dioceses — DE / DK / UK / SE / NO / CH) Copenhagen, Denmark · Remote from Berlin

- Shipped a capability into production under GDPR plus Protestant and Catholic data-protection law, across six countries.
- Mined support tickets and CRM records for recurring trends and feature requests, at a volume no one was going to read by hand.
- Product management and design together, one direct report, in a company of roughly 50.

Head of Product Design · Interim Head of Product

2023 – 2024

Ninox Software GmbH (no-code business platform — DACH / UK / US) Berlin, Germany

- Designed the AI onboarding model with the CEO and technical co-founder — prompt to database, with the schema growing turn by turn through suggested actions.
- Ran Claude over thousands of forum posts for sentiment and the recurring pain points underneath, with Dovetail as the shared repository.
- Practical experience developing AI agents across three companies, and the monitoring I have is qualitative rather than instrumented.
- Senior Leadership Team, five direct reports including a design lead. Hired two, promoted two.

Director of Design

2016 – 2021

Datameer GmbH (data analytics for top US / CA / UK banks and insurers) Berlin, Germany · team across Germany and NYC

- Five years of data-dense enterprise software under regulation, for banks and insurers in three countries.
- Functional head, two direct reports, peer among department heads under a VP of Design.

EARLIER CAREER PROGRESSION

Senior Mentor · *CareerFoundry* · 2016 – 2026 · Remote · a decade mentoring designers

Product & UX Strategist · *Hypoport hub SE* · 2021 – 2022 · Berlin · listed group, Holacracy

Earlier · *UX Strategy Lead, UX Manager, UX Designer, Web Developer* · 2009 – 2016 · Leipzig · Pennsylvania, USA

EDUCATION

M.A. Leadership

Kairos University, South Dakota, USA

Completed 2026

LANGUAGES

German · native

English · C2, fluent

SKILLS

Evaluating model output

A nine-criterion binary rubric, each judged in an isolated call · judge rationales returned to the generator · tiered criteria with declared evidence scope · judge-drift calibration · pre-registered A/B probes with byte-identical controls · blind preference judging · no golden set, deliberately

Prompt and knowledge engineering

A production agent’s system prompt, restricted-topic list and refusal rule · the grounding document authored and structured · quality criteria defined per message type before tuning · follow-up-question loops · prompt-to-database onboarding designed with a CEO and technical co-founder

Making systems refuse

A scope rule that declines to draft outside its remit · a suggestion engine bounded to the correct property family · returning nothing rather than something plausible and wrong · six automated checks identical locally and in CI

Unstructured text at volume

Thousands of forum posts analyzed with Claude for sentiment and recurring pain points · support tickets and CRM records mined for trends · interview transcripts transcribed and analyzed · Dovetail as the shared repository

Shipping under constraint

Production under GDPR alongside Protestant and Catholic data-protection law, six countries · agentic engineering daily · prototypes against production APIs that an external agency ships from · native German, C2 English