

Tobias Treppmann

Agentic AI Specialist — a production AI agent for a customer-service team, and the knowledge source behind it

LOCATION Berlin, Germany EMAIL hello@treppmann.design PORTFOLIO treppmann.design/for/docmorris LINKEDIN linkedin.com/in/tobiastreppmann

PROFILE

I built the evaluation panel that judges an AI pipeline's output, and shipped an AI agent for a customer-service team. Nine criteria, each judged separately, any failure rejecting the draft — no golden set. I wrote the agent's system prompt and author the knowledge source that decides what it can say. Designed and directed; implemented with AI coding agents.

EXPERIENCE

Head of Product Design

2025 – Present

Bibel TV (broadcaster in linear-to-streaming transformation) Hamburg + Berlin, Germany

- Shipped an AI support agent into production for the viewer-services team — wrote the system prompt, an explicit restricted-topic list, and the scope rule that refuses to draft outside a viewer enquiry.
- Author and structure the knowledge source it is grounded on — a web-hosted handbook summary. That page decides the agent's answerable scope, and is where most of the work actually goes.
- Designed, tested and optimized prompt structures to minimize hallucination risk and improve LLM performance, using grounded response logic and message-specific quality criteria.
- Rollout was voluntary — a self-test week, then one-to-ones. The skeptics came round after their hardest case, not the demo; every send stays a person's decision.
- Built an AI fundraising system end to end — analysis, prompt architecture, prototype, and a production Claude skill across in-app, web and video, with follow-up-question loops.
- Constrained a token suggestion engine after a numeric coincidence produced the wrong semantic relationship. It now returns nothing when no valid token qualifies.
- Six automated checks run locally and in CI, refusing unresolved tokens and hardcoded values before an external agency builds from the working prototype.
- Own UI quality across four platforms. Product managers prototype day to day, and an external engineering agency builds from my working reference.

Founder · Aperto

2025 – Present

Aperto (aperto.bible) — open Bible translation tooling and a multilingual content pipeline Berlin, Germany · independent

- Built a multi-layer evaluation panel for LLM output. Nine criteria, each judged in its own model call with no shared context, binary YES or NO, any single failure rejecting the candidate.
- No golden set, and that was a decision, not an omission. A reference key existed; removed in July 2026 after it hid eight of eleven failures on a chapter it had passed.
- Criteria are tiered and scoped: a correctness/mechanics floor, a register/reader-impact tier, and judge reasons returned verbatim inside a closed revision loop.
- Deterministic and statistical judges sit beside the model ones: style-distance, drift detection against published translations, and judge-drift calibration.
- Prompt changes validated by pre-registered A/B probes with byte-identical control arms, plus blind preference judging against the current baseline.
- Built the retrieval layer deliberately without embeddings — commentaries chunked by pericope into chapter-keyed JSON, so an exegesis step pulls exactly the scholarship for its passage.
- 139 loop runs and 3,138 individual criterion verdicts since July 2026, across 64 German, 55 English and 42 Polish chapters.
- Designed and directed across 1,384 commits; implemented with AI coding agents.

Head of Platform and Design

2024 – 2025

ChurchDesk (vertical ERP for churches and dioceses — DE / DK / UK / SE / NO / CH) Copenhagen, Denmark · Remote from Berlin

- Shipped a capability into production under GDPR plus Protestant and Catholic data-protection law, across six countries.
- Mined support tickets and CRM records for recurring trends and feature requests, at a volume no one was going to read by hand.
- Product management and design together, one direct report, in a company of roughly 50.

Head of Product Design · Interim Head of Product

2023 – 2024

Ninox Software GmbH (no-code business platform — DACH / UK / US) Berlin, Germany

- Designed the AI onboarding model with the CEO and technical co-founder — prompt to database, with the schema growing turn by turn through suggested actions.
- Ran Claude over thousands of forum posts for sentiment and the recurring pain points underneath, with Dovetail as the shared repository.
- Practical experience developing AI agents across three companies, and the monitoring I have is qualitative rather than instrumented.
- Senior Leadership Team, five direct reports including a design lead. Hired two, promoted two.

Director of Design

2016 – 2021

Datameer GmbH (data analytics for top US / CA / UK banks and insurers) Berlin, Germany · team across Germany and NYC

- Five years of data-dense enterprise software under regulation, for banks and insurers in three countries.
- Functional head, two direct reports, peer among department heads under a VP of Design.

EARLIER CAREER PROGRESSION

Senior Mentor · *CareerFoundry* · 2016 – 2026 · Remote · a decade mentoring designers

Product & UX Strategist · *Hypoport hub SE* · 2021 – 2022 · Berlin · listed group, Holacracy

Earlier · *UX Strategy Lead, UX Manager, UX Designer, Web Developer* · 2009 – 2016 · Leipzig · Pennsylvania, USA

EDUCATION

M.A. Leadership

Kairos University, South Dakota, USA

Completed 2026

LANGUAGES

German · native

English · C2, fluent

SKILLS

Evaluating model output

A nine-criterion binary rubric, each judged in an isolated call · judge rationales returned to the generator · tiered criteria with declared evidence scope · judge-drift calibration · pre-registered A/B probes with byte-identical controls · blind preference judging · no golden set, deliberately

Prompt and knowledge engineering

A production agent's system prompt, restricted-topic list and refusal rule · the grounding document authored and structured · hallucination-risk controls through grounding and scope · quality criteria defined per message type before tuning · follow-up-question loops · prompt-to-database onboarding designed with a CEO and technical co-founder

Making systems refuse

A scope rule that declines to draft outside its remit · a suggestion engine bounded to the correct property family · returning nothing rather than something plausible and wrong · six automated checks identical locally and in CI

Unstructured text at volume

Thousands of forum posts analyzed with Claude for sentiment and recurring pain points · support tickets and CRM records mined for trends · interview transcripts transcribed and analyzed · Dovetail as the shared repository

Shipping under constraint

Production under GDPR alongside Protestant and Catholic data-protection law, six countries · agentic engineering daily · prototypes against production APIs that an external agency ships from · native German, C2 English