

Tobias Treppmann

A production AI agent for a customer-service team, and the knowledge source behind it

CONTACT

LOCATION

Berlin, Germany

EMAIL

hello@treppmann.design

PORTFOLIO

treppmann.design/for/docmorris

LINKEDIN

linkedin.com/in/tobiastreppmann

EDUCATION

M.A. Leadership

Kairos University, South Dakota, USA
Completed 2026

LANGUAGES

German · native

English · C2, fluent

SKILLS

Evaluating model output

A nine-criterion binary rubric, each judged in an isolated call · judge rationales returned to the generator · tiered criteria with declared evidence scope · judge-drift calibration · pre-registered A/B probes with byte-identical controls · blind preference judging · no golden set, deliberately

Prompt and knowledge engineering

A production agent's system prompt, restricted-topic list and refusal rule · the grounding document authored and structured · quality criteria defined per message type before tuning · follow-up-question loops · prompt-to-database onboarding designed with a CEO and technical co-founder

Making systems refuse

A scope rule that declines to draft outside its remit · a suggestion engine bounded to the correct property family · returning nothing rather than something plausible and wrong · six automated checks identical locally and in CI

Unstructured text at volume

Thousands of forum posts analyzed with Claude for sentiment and recurring pain points · support tickets and CRM records mined for trends · interview transcripts transcribed and analyzed · Dovetail as the shared repository

Shipping under constraint

Production under GDPR alongside Protestant and Catholic data-protection law, six countries · agentic engineering daily · prototypes against production APIs that an external agency ships from · native German, C2 English

PROFILE

I built the evaluation panel that judges an AI pipeline's output, and I shipped an AI agent into production for a customer-service team. The panel runs nine criteria, each judged in its own model call, YES or NO, any failure rejecting the candidate and the reasons fed back to the generator. There is no golden set on purpose: the reference key I started with was removed once it turned out to be hiding failures. On the agent side I wrote the system prompt, the restricted topics and the refusal rule, and I author the knowledge source it is grounded on. Designed and directed; implemented with AI coding agents.

3,138 verdicts

individual criterion judgments across
139 evaluation loop runs since July

I removed my own key

the reference set was hiding 8 of 11
failures on a chapter it passed

EXPERIENCE

Head of Product Design

2025 – Present

Bibel TV (broadcaster in linear-to-streaming transformation) · Hamburg + Berlin, Germany

- Shipped an AI support agent into production for the viewer-services team, drafting replies to viewer letters in the service handbook's wording.
- The system prompt is mine, with an explicit restricted-topic list and a scope rule that refuses to draft when the input is not a viewer enquiry.
- Created and structured the conversation flow and response logic: what the agent answers, what it declines, and when it asks for clarification instead of guessing.
- I author and structure the knowledge source it is grounded on — a web-hosted handbook summary carrying snippets, standard phrasing and the service team's own wording.
- That page decides the agent's answerable scope, which is the part of the work I spend most time on. Optimizing the interaction starts there, not in the prompt.
- Minimizing hallucination started before the prompt. We defined with the support team what a good draft looks like per message type, then tested and optimized against that.
- Rollout was voluntary. A self-test week, then one-to-ones, encouraged by team leads rather than mandated.
- The sceptics came round after their hardest case, not after the demo. Every send is still a person's decision.
- Built an AI fundraising system end to end — analysis, prompt architecture, prototype, and a production Claude skill across in-app, web and video.
- It runs follow-up-question loops, so the model asks before it generates rather than filling gaps by inference.
- Built a token suggestion engine, then constrained it. It offered a spacing token as the fix for a 12px font size because the numbers matched.
- It would have rendered pixel-identical and welded body text to a gap between cards. Suggestions are bounded to the property's own family now.
- Where no token qualifies, the check says nothing at all. A confidently wrong suggestion is worse than none for someone who does not know the catalog.
- Enforcement runs as six checks, one command, identical locally and in CI, refusing an unresolvable token reference or a value written straight into a component.
- Prototype in production code against real APIs, and an external engineering agency ships from that code.
- Claude Code has been the daily tool for over a year.
- Design, test, adjust, test again. Prompt structures are optimized against written criteria, not against a feeling.
- It refuses. That was the design.
- One agent, in production, in German.

Founder · Aperto

2025 – Present

Aperto (aperto.bible) — open Bible translation tooling and a multilingual content pipeline · Berlin, Germany · independent

- Built a multi-layer evaluation panel for LLM output. Nine criteria, each judged in its own model call with no shared context, binary YES or NO, and any single failure rejects the candidate.
- No golden set, and that was a decision rather than an omission. A reference key existed and I removed it in July 2026 after finding it had hidden eight of eleven failures on a chapter it passed.
- Criteria are tiered and scoped. A floor tier covers correctness and mechanics; the primary tier covers register and reader impact. Each criterion declares what the judge is allowed to see.
- Judge reasons go back to the generator verbatim, inside a closed loop: judge, apply findings, validate, re-judge, halt on convergence stall.
- The loop revises the failing span rather than re-rolling the whole draft, and reuses verdicts for criteria whose spans were untouched.
- Deterministic and statistical judges sit beside the model ones. Statistical style-distance against a reference-author corpus, using KL divergence over sentence length, lexical density and type-token ratio.

- It is documented as a floor, not a verdict. A number that cannot decide is worse than no number if anyone treats it as one.
- Drift detection against three published translations per language, by token-set Jaccard and character trigram Dice. Agreeing with all three is as interesting as agreeing with none.
- A Judge-drift control re-judges a flagged criterion against a calibration anchor, because a judge that moves is worse than a judge that is wrong.
- Prompt changes are validated by pre-registered A/B probes with the control arm verified byte-identical, plus blind preference judging against the current baseline.
- Built the retrieval layer deliberately without embeddings. Commentaries extracted from PDF, EPUB, text and markdown into JSON, chunked by pericope and keyed by book and chapter.
- Structural retrieval, so an exegesis step always pulls exactly the scholarship for its passage. Sentence-transformers were rejected for determinism and the reasoning is in the code.
- 139 loop runs and 3,138 individual criterion verdicts since July 2026, across a corpus of 64 German, 55 English and 42 Polish chapters.
- Designed and directed by me across 1,384 commits. Implemented with AI coding agents.
- Nine criteria. Nine separate calls.
- Any single NO rejects the draft. No partial credit.

Head of Platform and Design

2024 – 2025

ChurchDesk (vertical ERP for churches and dioceses — DE / DK / UK / SE / NO / CH) · Copenhagen, Denmark · Remote from Berlin

- Shipped a capability into production under GDPR plus Protestant and Catholic data-protection law, across six countries.
- Mined support tickets and CRM records for recurring trends and feature requests, at a volume no one was going to read by hand.

Head of Product Design · Interim Head of Product

2023 – 2024

Ninox Software GmbH (no-code business platform — DACH / UK / US) · Berlin, Germany

- Designed the AI onboarding model with the CEO and technical co-founder — prompt to database, with the schema growing turn by turn through suggested actions.
- Ran Claude over thousands of forum posts for sentiment and the recurring pain points underneath, with Dovetail as the shared repository.
- Practical experience developing AI agents across three companies, and the monitoring I have is qualitative rather than instrumented.

Director of Design

2016 – 2021

Datameer GmbH (data analytics for top US / CA / UK banks and insurers) · Berlin, Germany · team across Germany and NYC

- Five years of data-dense enterprise software under regulation, for banks and insurers in three countries.

EARLIER CAREER PROGRESSION

Senior Mentor · *CareerFoundry* · 2016 – 2026 · Remote · a decade mentoring designers

Product & UX Strategist · *Hypoport hub SE* · 2021 – 2022 · Berlin · listed group, Holacracy

Earlier · *UX Strategy Lead, UX Manager, UX Designer, Web Developer* · 2009 – 2016 · Leipzig · Pennsylvania, USA